

Four use cases. One Kaltura stack. Grounded in Admiral's own approved content.

A working summary of what was demonstrated, prepared for Amber, Adam, Daniel and Nadine to use as the basis for Admiral's internal priority-setting conversation.

The frame. Each use case stands alone, but the data, the agents and the analytics flow between them. The decision isn't *which* use cases — it's *which one starts first*, knowing the others come along behind it on the same stack.

USE CASE 1

Role-play training for contact-centre agents

TODAY

1 trainer per ~15 learners. In peer-to-peer role-play, agents give each other the easiest version of the scenario — useful for confidence, less useful for realism.

KALTURA

Avatar plays the customer. Configurable tone, difficulty, emotional arc. Generative — when the agent answers poorly, the avatar drills deeper into *that* area in real time. Custom scoring dimensions, multiple scorecard templates (manager vs front-line). Results sync to Cornerstone / LRS.

USE CASE 2 · THE WEDGE

Screen-share assessment

TODAY — IN DANIEL'S WORDS

"It knows what a good call sounds like and it can transcribe and it'll read like a really good call. But we've just got no evidence that what has happened on that call has happened on the system." Verify: price quoted vs price set on policy; products requested vs products added; driver / cover changes correctly applied.

KALTURA

Avatar watches the agent's screen, compares actions in the system against the call transcript (or any ground-truth source), validates field-by-field, triggers human review when accuracy drops below threshold.

USE CASE 3

Learning concierge across knowledge sites

TODAY

Agents hunt across help hubs, Cornerstone, scripts, process docs. Keyword search returns the wrong answer. Delhi agents face the additional language barrier.

KALTURA

Single conversational layer over Admiral's approved content. Answers in plain English (or any language), cites the source, logs every interaction. Patterns of struggle drive Training Needs Analysis — the avatar surfaces topics where multiple agents need help, before TNA has to find them manually.

USE CASE 4

Generative content + assessment

TODAY — AMBER'S STATED NEED

"We'd love a solution for how we test competency aside from role-plays — quizzes or regular assessments. Currently using automation, but it's clunky."

KALTURA

Given a topic and a configured persona, the avatar generatively builds quiz questions, walks the learner through, marks the answers, produces a summary file mappable to CRM or LRS fields. Switches between teacher mode and quiz mode in the same session.

HOW THE FOUR STACK TOGETHER

Admiral approved content — Cornerstone · LRS · scripts · process docs · policy data · call recordings · service-link transcripts → **Kaltura avatar stack** — persona engine · constraint layer · conversational AI · screen understanding · scoring engine · audit log → Role-play · Screen-share assessment · Concierge · Generative content. One stack, one knowledge base, one audit trail, four surfaces.

WHAT'S WORTH DECIDING INTERNALLY BEFORE THE ON-SITE

1 Which use case starts first? Role-play = lowest integration overhead. Screen-share = highest visible value. Concierge = broadest user base. Quizzes = lowest user friction.

2 Which cohort first? Induction (Adam's area) is a clean, contained pilot. Continuing development reaches the 3,000 immediately.

3 Where does it land first? Cornerstone-embedded, LRS-fed, or a parallel surface? Affects integration scope, not the avatar build.

4 What anchor data can Admiral share in advance? Anything we can ingest before the on-site lets us show an Admiral-specific demo, not a generic one.

Derek Grant · Enterprise Sales Lead, European Financial Services · Kaltura
Ben English · Director of Solutions Engineering · Kaltura

Next: Kaltura UK on-site at Swansea Park, second half of June 2026.
This is a working summary, not a proposal.